

Topic 5: Creating and flattening Docker images

Dr. Daniel Yellin

Building image from images, flattening images

THESE SLIDES ARE THE PROPERTY OF DANIEL YELLIN
THEY ARE ONLY FOR USE BY STUDENTS OF THE CLASS
THERE IS NO PERMISSION TO DISTRIBUTE OR POST
THESE SLIDES TO OTHERS


Overview

- Docker review
- Producing images from containers
- Making your container skinny

Docker review

Images


An *image* is a collection of files, metadata, and a command. It contains executable application source code as well as all the tools, libraries, and dependencies that the application code needs to run as a container.



Containers

A *container* is created from an image.


- It is an isolated running process (or processes) started using the command, with access only to the files in the image. The metadata defines properties of the process such as network ports accessible to the process.
- You can create, start, stop, move, or delete a container using the Docker API or CLI. You can connect a container to one or more networks, attach storage to it, or even create a new image based on its current state.



Dockerfiles

A *Dockerfile* automates the process of Docker image creation. It's a list of command-line interface (CLI) instructions that Docker Engine will run in order to assemble the image.

Often, an image is based upon another image. To build your own image, you create a Dockerfile with a simple syntax for defining the steps needed to create your desired image “on top of” the other image.



Docker image layers example

This Docker file results in several layers:

- A layer for the base image, `python:alpine3.17` (this itself is comprised of layers).
- A layer with the `toys.py` files
- A layer with the Flask files installed

The other commands, such as “`WORKDIR`”, is metadata that does not modify the filesystem and is therefore not part of the image.



```
FROM python:alpine3.17
WORKDIR ./app
COPY toys.py .
RUN pip install Flask
ENV FLASK_APP=toys.py
ENV FLASK_RUN_PORT=8001
#ENV NINJA_API_KEY - if required
EXPOSE 8001
CMD ["flask", "run", "--host=0.0.0.0"]
```

Image layers

- **Read-only and Immutable:** Once a layer is created, it becomes read-only and cannot be modified.
- **Layered Construction:** Each instruction in a Dockerfile that modifies the filesystem (e.g., RUN, COPY, ADD) creates a new layer.

These layers are stacked sequentially, with the base image forming the bottom layer and subsequent instructions adding layers on top. Any Dockerfile instruction that modifies the file system creates a new layer.

Layer caching and Copy-on-Write

- **Layer Caching:** Docker utilizes layer caching during the image build process. If a layer and its preceding layers remain unchanged from a previous build, Docker can reuse the existing cached layer, significantly speeding up subsequent builds.
 - This is particularly beneficial when only application code changes, as dependency installation layers can often be reused.
- **Copy-on-Write (CoW):** Docker employs a copy-on-write strategy for layers. When a container modifies a file that exists in a lower, read-only layer, a copy of that file is made into the container's writable layer, and the modification is applied to the copy. The original file in the lower layer remains untouched.

The “-it” flags in the Docker run command

docker run <image>

Creates and runs a new container from the given image. You cannot directly interact with the container from your shell.

docker run -it <image>

Creates and runs a new container from the given image. The “-it” flags (interactive, terminal device) allows you to interact with the image via the shell.

The Docker exec command

docker exec -it <container> <cmd>

Executes the given command in the specified running container.

Example: First I start up a toys container via the cmd **docker run toys**

At this point, it is running but I cannot “get inside” the container.

```
danielyellin@Daniels-Air ~ % docker ps
CONTAINER ID   IMAGE    COMMAND           CREATED          STATUS          PORTS     NAMES
33b5ecf25d61   toys    "flask run --host=0..."   12 seconds ago   Up 11 seconds   8001/tcp   cool_mccarthy
danielyellin@Daniels-Air ~ % docker exec -it cool_mccarthy sh
/app # ls
__pycache__  toys.py
/app #
```

Runs the sh cmd inside the container

Now I have an interactive shell and I can issue cmds inside the container (like ls).

Producing images from containers

How to make an image without a Dockerfile

Container state

We learnt previously how to make a container from an image. In this section we will show how to make an image from a container!

Recall that:

- Everytime you create and run a container it has the *same filesystem* as the image it is created from.
- If a container makes changes to its file system and then the image is *stopped*, the changes to the file system are preserved when the container is *started* up again. However, when the image is removed (killed), the changes it made are lost.

Committing a container to yield a new image

1. Run a container (built on a base image)
2. Run some commands inside the container changing the file system
3. **Commit** the container

The result is a new image that contains the changes you made.

Let's see an example

Making an image via container commit

Start off by the command

```
1. docker pull  
alpine:latest
```

This retrieves the latest alpine image from Docker Hub

```
[danielyellin@Daniels-MacBook-Air stocks % docker pull alpine  
Using default tag: latest  
latest: Pulling from library/alpine  
9986a736f7d3: Pull complete  
Digest: sha256:1e42bbe2508154c9126d48c2b8a75420c3544343bf86fd  
Status: Downloaded newer image for alpine:latest  
docker.io/library/alpine:latest
```

[repo-info](#) repo's [repos/alpine/](#) directory (history)

(image metadata, transfer size, etc)

- **Image updates:**

[official-images](#) repo's [library/alpine](#) label

[official-images](#) repo's [library/alpine](#) file (history)

- **Source of this description:**

[docs](#) repo's [alpine/](#) directory (history)

What is Alpine Linux?

Alpine Linux is a Linux distribution built around [musl](#) [libc](#) and [BusyBox](#). The image is only 5 MB in size and has access to a [package repository](#) that is much more complete than other BusyBox based images. This makes Alpine Linux a great image base for utilities and even production applications. [Read more about Alpine Linux here](#) and you can see how their mantra fits in right at home with Docker images.



How to use this image

Usage

Use like you would any other base image:

```
FROM alpine:3.14  
RUN apk add --no-cache mysql-client  
ENTRYPOINT ["mysql"]
```

Making an image via container commit (cont)

2. docker run -it alpine:latest /bin/sh

This creates a container running the alpine image. The flag “-it” makes the container an interactive session connected to the terminal window. “/bin/sh” tells it to run the shell when it starts up. We issue the following cmd **inside** the container.

3. echo “I love docker” > my_file.txt

This creates a file named “my_file.txt” in the container. The file contains the string “I love Docker”

```
danielyellin@Daniels-Air topic3-docker % docker run -it alpine:latest /bin/sh
/ # ls
bin dev etc home lib media mnt opt proc root run sbin srv sys tmp usr var
/ #
/ # echo "I love docker" > my_file.txt
/ # ls
bin etc lib mnt opt root sbin sys usr
dev home media my_file.txt proc run srv tmp var
/ # exit
```

Making an image via container commit (cont)

4 . Exit the container. This stops the container but it still exists.

5 . docker ps -a

We see that since we did not specify a name for this container, Docker gave it the name “eager_carver”

```
danielyellin@Daniels-Air topic3-docker % docker ps -a
CONTAINER ID        IMAGE           COMMAND       CREATED          STATUS          PORTS          NAMES
d28087fa446e        alpine:latest   "/bin/sh"     39 seconds ago  Exited (0) 5 seconds ago
0eab87ec3d8e        alpine          "/bin/sh"     21 minutes ago  Exited (0) 20 minutes ago
d580e6949a21        alpine:latest   "-it /bin/sh"  25 minutes ago  Created
5d4de57425a4        alpine:latest   "/bin/sh"     25 minutes ago  Exited (0) 25 minutes ago
ff4f3ee81404        restsvr-v1    "python3 rest-word-s..." 3 days ago      Exited (0) 55 minutes ago
danielyellin@Daniels-Air topic3-docker % docker commit eager_carver my_new_image
sha256:86c523767517656d00caad36e87d0c257964aefe29b0f5bd56ee2702d638d375
danielyellin@Daniels-Air topic3-docker % docker images
REPOSITORY          TAG           IMAGE ID        CREATED          SIZE
my_new_image        latest         86c523767517  11 seconds ago  7.46MB
restsvr-v1          latest         043bb415df9e  3 days ago      75.8MB
alpine              latest         04eeaa5f8c35  13 days ago      7.46MB
```

Making an image via container commit (cont)

6. `docker commit eager_carver my_new_image`

This cmd makes a new image from the container “eager_carver” and we give this image the name “my_new_image”. It stores the state of the file system (not the process).

7. `docker run -it my_new_image /bin/sh`

This starts up a container using the new image we just created - the alpine image plus a modified file system (a new file `my_file.txt`).

The `-it` flag runs the image iteractively and connects it to your terminal window

Making an image via container commit (cont)

8. Inside the running container issue the command ls
9. Then issue the command cat myfile.txt


We see that the file we created is in the image. We successfully modified an existing image and created a new one via the commit command.

```
[danielyellin@Daniels-Air topic3-docker % docker run -it my_new_image /bin/sh
[ / # ls
bin          etc          lib          mnt          opt          root         sbin         sys          usr
dev          home         media        my_file.txt  proc         run          srv          tmp          var
[ / # cat my_file.txt
I love docker
[ / # exit
```

Images and containers

An *image* is created from a container

- By committing a container, you get an image.
- And that is how Dockerfile actually builds containers
 1. It builds a container with the base image
 2. It executes the next Dockerfile cmd, and commits to create a new image
 3. It uses that new image to build a new container
 4. It repeats steps 2-3 until there are no more commands




Making your container skinny

Flattening a container - How to get rid of layers


Summary: Docker commit

We just saw that you can create a new image by modifying the file system of a container and then *committing* that container. It becomes a new image.



Example

1. Initial image: alpine
2. Run the alpine image in a container and add a file secret.txt
3. Commit that container to an image called layer1
4. Run the layer1 image in a container, add a file secret2.txt and delete secret.txt
5. Commit that container to an image called layer2



What is the difference in the file systems in the two images?

Example (cont)

- The file system in layer1 looks like this:

```
[/ # ls
bin          etc          lib          mnt          proc         run
dev          home         media        opt          root         sbin
secret.txt   srv          secret.txt  sy          tr]
```

- The file system in layer2 looks like this:

```
[/ # echo "another secret" > secret2.txt
[ / # ls
bin          etc          lib          mnt          proc         run
dev          home         media        opt          root         sbin
secret.txt   secret2.txt  secret.txt  sr          sy
[ / # rm secret.txt
[ / # ls
bin          etc          lib          mnt          proc         run
dev          home         media        opt          root         sbin
secret2.txt  srv          secret2.txt sy          tm]
```

- What do you think the different image layers will be in these images?*

Example (cont)

- docker history layer1

```
[danielyellin@Daniels-Air ~ % docker history layer1
IMAGE          CREATED      CREATED BY
71bcba8a6a9c  14 minutes ago /bin/sh
d74e625d9115  5 weeks ago   /bin/sh -c #(nop)  CMD ["/bin/sh"]
<missing>      5 weeks ago   /bin/sh -c #(nop) ADD file:9bd9ea42a9f3bdc76...
                                         SIZE      COMMENT
                                         48B      Added secret.txt
                                         0B       Start cmd
                                         7.46MB  Initial image
```


- docker history layer2

```
[danielyellin@Daniels-Air ~ % docker history layer2
IMAGE          CREATED      CREATED BY
0d9fba4bafda  9 seconds ago /bin/sh
71bcba8a6a9c  55 minutes ago /bin/sh
d74e625d9115  5 weeks ago   /bin/sh -c #(nop)  CMD ["/bin/sh"]
<missing>      5 weeks ago   /bin/sh -c #(nop) ADD file:9bd9ea42a9f3bdc76...
                                         SIZE      COMMENT
                                         117B     Added secret2.txt &
                                         48B      deleted secret.txt
                                         0B       Added secret.txt
                                         7.46MB  Start cmd
                                         Initial image
```

Example (cont)

- Next remove the containers and delete all the images exempt *layer2*

- *QUESTION:* can we recover *secret.txt*?



Example (cont)

- Recall the history of layer2, all the layers are available in the image

IMAGE	CREATED	CREATED BY	SIZE	COMMENT
0d9fba4bcfda	9 seconds ago	/bin/sh	117B	Added secret2.txt & deleted secret.txt
71bcba8a6a9c	55 minutes ago	/bin/sh	48B	Added secret.txt
d4e625d9115	5 weeks ago	/bin/sh - c #(nop) CMD ["/bin/sh"]	0B	Start cmd
<missing>	5 weeks ago	/bin/sh - c #(nop) ADD file:9bd9ea42a9f3bdc76...	7.46MB	Initial image


Lets see what happens when we issue the command:

```
docker run -it 71bcba8a6a9c /bin/sh
```

Example (cont)

- Because layer2 is on top of layer1, we can “recover” the filesystem snapshot from the time of layer1 !


```
[danielyellin@Daniels-Air ~ % docker run -it 71bcba8a6a9c /bin/sh
/d # ls
bin      etc      lib      mnt      proc      run      secret.txt  sys      us
dev      home     media    opt      root      sbin     srv      tmp      va
/d # cat secret.txt
my secret
/d # ]
```



Example (cont)

- We often do not want this behavior because:
 - It exposes secrets we do not want others to see
 - It blows up the size of our image (if there is layer that has a lot of large files that is no longer needed)
- There are multiple ways to *flatten* an image
- One way is to use **docker export** to create a *tar file*, then to use **docker import** to create an image from that tar file. When you do this, all the intermediate images are removed.

Docker export and import



Example (cont)

```
[danielyellin@Daniels-Air ~ % docker run -d layer2 /bin/true
1d6deaa8ba5cdc8e9ec33e1f633d100d4d0db4f232ec285d249d988c48be16cd
[danielyellin@Daniels-Air ~ % docker export 1d6de | docker import - flat
sha256:1a513e2d0f098f8f8c299e83d1653ebf1604df4b814e386f4dfdc20605baed14
[danielyellin@Daniels-Air ~ % docker images
REPOSITORY          TAG      IMAGE ID      CREATED       SIZE
flat                latest   1a513e2d0f09  7 seconds ago  7.46MB
```

Creates a container from layer2 image and exits
This is the ID of the container running layer2
See explanation of this cmd below

~~docker export 1d6de | docker import - flat~~

- The first part of this line is a command that exports the image in the container beginning with ID 1d6de (layer2 image)
- It then *pipes* the output of this command, a *tar* file, into the next command
- The last part of this line is a that command takes a tar file and outputs an image of the given name; in our case we call this image “flat”
- This is the same as issuing the commands:
 - `docker export 1d6de my_image_tar`
 - `docker import my_image_tar flat`

Example (cont)

- Let's look at the history of the flattened image. Let's also run the flattened image and look inside at the file system.


```
[danielyellin@Daniels-Air ~ % docker history flat
IMAGE          CREATED      CREATED BY      SIZE      COMMENT
1a513e2d0f09  51 seconds ago    7.46MB      Imported from -
[danielyellin@Daniels-Air ~ % docker run -it flat /bin/sh
[/ # ls
bin          etc          lib          mnt          proc        run        secret2.txt  s
dev          home         media        opt          root       sbin      srv          t
[/ # exit
```

Only a single layer
is visible in the
history

Only secret2.txt,
not secret.txt

- Run the history cmd and you will not see any additional layers – they are gone.

Flattening images



Reasons to flatten images

- **Reduced size** – Flattening an image can significantly reduce its size, as it removes the overhead of multiple layers.
- **Improved performance** – Flattening an image can improve the performance of the image, as it reduces the number of layers that must be processed during runtime.
- **Simplified distribution** – Flattening an image can simplify the process of distributing the image, as it removes the need to manage multiple layers and their dependencies.
- **Remove “secrets”** – Flattening an image removes information that might be in early layers of the image, but not intended for distribution.

Ordering of Dockerfile commands to reduce image size

- Docker caches images, so when you build a Docker image from a buildfile, it will look at each step and see if an existing image up to that step exists in the cache.
- If changes cause a step to be redone (e.g., a file it relies on has changes), then a new image is built, and all the steps forward will create new images.
- For Dockerfiles with steps that require a lot of time to execute, you want to order commands that may change Docker image state to be as late as possible in the Dockerfile.

For example, if you place an ADD instruction towards the top of the Dockerfile and anything in the added file or directory has changed, then all subsequent layers associated with commands below the ADD will have their cached images invalidated.

[Source1](#) and [Source2](#)